

Importing the Dependencies

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from xgboost import XGBRegressor
from sklearn import metrics
```

Data Collection & Processing

```
# loading the data from csv file to a Pandas DataFrame
calories = pd.read_csv('/content/calories.csv')

# print the first 5 rows of the dataframe
calories.head()
```

|   | User_ID  | Calories |
|---|----------|----------|
| 0 | 14733363 | 231.0    |
| 1 | 14861698 | 66.0     |
| 2 | 11179863 | 26.0     |
| 3 | 16180408 | 71.0     |
| 4 | 17771927 | 35.0     |

```
exercise_data = pd.read_csv('/content/exercise.csv')

exercise_data.head()
```

|   | User_ID  | Gender | Age | Height | Weight | Duration | Heart_Rate | Body_Temp |
|---|----------|--------|-----|--------|--------|----------|------------|-----------|
| 0 | 14733363 | male   | 68  | 190.0  | 94.0   | 29.0     | 105.0      | 40.8      |
| 1 | 14861698 | female | 20  | 166.0  | 60.0   | 14.0     | 94.0       | 40.3      |
| 2 | 11179863 | male   | 69  | 179.0  | 79.0   | 5.0      | 88.0       | 38.7      |
| 3 | 16180408 | female | 34  | 179.0  | 71.0   | 13.0     | 100.0      | 40.5      |
| 4 | 17771927 | female | 27  | 154.0  | 58.0   | 10.0     | 81.0       | 39.8      |

Combining the two Dataframes

```
calories_data = pd.concat([exercise_data, calories['Calories']], axis=1)

calories_data.head()
```

|   | User_ID  | Gender | Age | Height | Weight | Duration | Heart_Rate | Body_Temp | Calories |
|---|----------|--------|-----|--------|--------|----------|------------|-----------|----------|
| 0 | 14733363 | male   | 68  | 190.0  | 94.0   | 29.0     | 105.0      | 40.8      | 231.0    |
| 1 | 14861698 | female | 20  | 166.0  | 60.0   | 14.0     | 94.0       | 40.3      | 66.0     |
| 2 | 11179863 | male   | 69  | 179.0  | 79.0   | 5.0      | 88.0       | 38.7      | 26.0     |
| 3 | 16180408 | female | 34  | 179.0  | 71.0   | 13.0     | 100.0      | 40.5      | 71.0     |
| 4 | 17771927 | female | 27  | 154.0  | 58.0   | 10.0     | 81.0       | 39.8      | 35.0     |

```
# checking the number of rows and columns
calories_data.shape

(15000, 9)

# getting some informations about the data
calories_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 15000 entries, 0 to 14999
Data columns (total 9 columns):
#   Column      Non-Null Count  Dtype
---  -
0   User_ID     15000 non-null  int64
```

```
1  Gender      15000 non-null object
2  Age         15000 non-null int64
3  Height      15000 non-null float64
4  Weight      15000 non-null float64
5  Duration    15000 non-null float64
6  Heart_Rate  15000 non-null float64
7  Body_Temp   15000 non-null float64
8  Calories    15000 non-null float64
dtypes: float64(6), int64(2), object(1)
memory usage: 1.0+ MB
```

```
# checking for missing values
calories_data.isnull().sum()
```

```
User_ID      0
Gender        0
Age           0
Height        0
Weight        0
Duration      0
Heart_Rate    0
Body_Temp     0
Calories      0
dtype: int64
```

Data Analysis

```
# get some statistical measures about the data
calories_data.describe()
```


|       | User_ID      | Age          | Height       | Weight       | Duration     | Heart_Rate   | Body_Temp    | Calories     |
|-------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|
| count | 1.500000e+04 | 15000.000000 | 15000.000000 | 15000.000000 | 15000.000000 | 15000.000000 | 15000.000000 | 15000.000000 |
| mean  | 1.497736e+07 | 42.789800    | 174.465133   | 74.966867    | 15.530600    | 95.518533    | 40.025453    | 89.539533    |
| std   | 2.872851e+06 | 16.980264    | 14.258114    | 15.035657    | 8.319203     | 9.583328     | 0.779230     | 62.456978    |
| min   | 1.000116e+07 | 20.000000    | 123.000000   | 36.000000    | 1.000000     | 67.000000    | 37.100000    | 1.000000     |
| 25%   | 1.247419e+07 | 28.000000    | 164.000000   | 63.000000    | 8.000000     | 88.000000    | 39.600000    | 35.000000    |
| 50%   | 1.499728e+07 | 39.000000    | 175.000000   | 74.000000    | 16.000000    | 96.000000    | 40.200000    | 79.000000    |
| 75%   | 1.744928e+07 | 56.000000    | 185.000000   | 87.000000    | 23.000000    | 103.000000   | 40.600000    | 138.000000   |
| max   | 1.999965e+07 | 79.000000    | 222.000000   | 132.000000   | 30.000000    | 128.000000   | 41.500000    | 314.000000   |

Data Visualization

```
sns.set()
```


```
# plotting the gender column in count plot
sns.countplot(calories_data['Gender'])
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. Fr
FutureWarning
<matplotlib.axes._subplots.AxesSubplot at 0x7fcbbd756110>
```




```
# finding the distribution of "Age" column
sns.distplot(calories_data['Age'])
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be
warnings.warn(msg, FutureWarning)
<matplotlib.axes._subplots.AxesSubplot at 0x7fcbbd200550>
```




```
# finding the distribution of "Height" column
sns.distplot(calories_data['Height'])
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be
warnings.warn(msg, FutureWarning)
<matplotlib.axes._subplots.AxesSubplot at 0x7fcbb1ed3d10>
```



```
# finding the distribution of "Weight" column
sns.distplot(calories_data['Weight'])
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be
warnings.warn(msg, FutureWarning)
<matplotlib.axes._subplots.AxesSubplot at 0x7fcbb1e2c190>
```




### Finding the Correlation in the dataset

1. Positive Correlation
2. Negative Correlation

```
correlation = calories_data.corr()
```

```
# constructing a heatmap to understand the correlation
```

```
plt.figure(figsize=(10,10))
sns.heatmap(correlation, cbar=True, square=True, fmt='.1f', annot=True, annot_kws={'size':8}, cmap='Blues')
```



Converting the text data to numerical values

```
calories_data.replace({"Gender":{"male":0,'female':1}}, inplace=True)

calories_data.head()
```

|   | User_ID  | Gender | Age | Height | Weight | Duration | Heart_Rate | Body_Temp | Calories |
|---|----------|--------|-----|--------|--------|----------|------------|-----------|----------|
| 0 | 14733363 | 0      | 68  | 190.0  | 94.0   | 29.0     | 105.0      | 40.8      | 231.0    |
| 1 | 14861698 | 1      | 20  | 166.0  | 60.0   | 14.0     | 94.0       | 40.3      | 66.0     |
| 2 | 11179863 | 0      | 69  | 179.0  | 79.0   | 5.0      | 88.0       | 38.7      | 26.0     |
| 3 | 16180408 | 1      | 34  | 179.0  | 71.0   | 13.0     | 100.0      | 40.5      | 71.0     |
| 4 | 17771927 | 1      | 27  | 154.0  | 58.0   | 10.0     | 81.0       | 39.8      | 35.0     |

Separating features and Target

```
X = calories_data.drop(columns=['User_ID','Calories'], axis=1)
Y = calories_data['Calories']

print(X)

   Gender  Age  Height  Weight  Duration  Heart_Rate  Body_Temp
0       0   68   190.0    94.0     29.0     105.0     40.8
1       1   20   166.0    60.0     14.0     94.0     40.3
2       0   69   179.0    79.0      5.0     88.0     38.7
3       1   34   179.0    71.0     13.0    100.0     40.5
4       1   27   154.0    58.0     10.0     81.0     39.8
...     ...  ...    ...    ...     ...     ...     ...
14995    1   20   193.0    86.0     11.0     92.0     40.4
14996    1   27   165.0    65.0      6.0     85.0     39.2
14997    1   43   159.0    58.0     16.0     90.0     40.1
14998    0   78   193.0    97.0      2.0     84.0     38.3
14999    0   63   173.0    79.0     18.0     92.0     40.5

[15000 rows x 7 columns]

print(Y)

0      231.0
1       66.0
2       26.0
3       71.0
4       35.0
```

```
...
14995    45.0
14996    23.0
14997    75.0
14998    11.0
14999    98.0
Name: Calories, Length: 15000, dtype: float64
```

Splitting the data into training data and Test data

```
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_state=2)
```

```
print(X.shape, X_train.shape, X_test.shape)
```

```
(15000, 7) (12000, 7) (3000, 7)
```

Model Training

XGBoost Regressor

```
# loading the model
model = XGBRegressor()
```

```
# training the model with X_train
model.fit(X_train, Y_train)
```

```
[10:06:32] WARNING: /workspace/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of reg:squarederror.
XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,
             colsample_bynode=1, colsample_bytree=1, gamma=0,
             importance_type='gain', learning_rate=0.1, max_delta_step=0,
             max_depth=3, min_child_weight=1, missing=None, n_estimators=100,
             n_jobs=1, nthread=None, objective='reg:linear', random_state=0,
             reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
             silent=None, subsample=1, verbosity=1)
```

Evaluation

Prediction on Test Data

```
test_data_prediction = model.predict(X_test)
```

```
print(test_data_prediction)
```

```
[129.06204  223.79721  39.181965 ... 145.59767  22.53474  92.29064 ]
```

Mean Absolute Error

```
mae = metrics.mean_absolute_error(Y_test, test_data_prediction)
```

```
print("Mean Absolute Error = ", mae)
```

```
Mean Absolute Error =  2.7159012502233186
```